

On the CUDA implementation of Reduce and Scan

Cosmin Oancea and Troels Henriksen

September 2025

Course Contents

Implementation of Reduce


Implementation of Scan

Summing an Integer Array

$$\sum_{i < n} x[i]$$

Binary Tree Reduction

The idea: each thread reads two neighbouring elements, adds them together, and writes one element. This halves the array in size. Continue until only a single element is left.



- ▶ Each level becomes a kernel invocation, with number of threads equal to half the number of array elements.
- ▶ $O(n)$ work and $O(\log(n))$ span (optimal).
- ▶ **Why is this not efficient?**

Improving the Tree Reduction

The idea: instead of shrinking the array by a factor of two for each level, shrink it by the CUDA-block size.

- ▶ Same asymptotic performance.
- ▶ Avoids kernels with very few threads. E.g with block size 256:
 $10000000 \rightarrow 39063 \rightarrow 153 \rightarrow 1$.

Improving the Tree Reduction

The idea: instead of shrinking the array by a factor of two for each level, shrink it by the CUDA-block size.

- ▶ Same asymptotic performance.
- ▶ Avoids kernels with very few threads. E.g with block size 256:
 $10000000 \rightarrow 39063 \rightarrow 153 \rightarrow 1$.

Implementation	$n = 1000$	$n = 1000000$
Tree reduction	$77\mu s$	$363\mu s$
Block reduction	$17\mu s$	$179\mu s$

Applying Brent's Lemma

The idea: instead of letting the thread count depend on the input size, always launch the same number of threads, and have each thread perform an efficient sequential summation of a *chunk* of the input.

- ▶ GPUs have a maximum (hardware/problem-dependent) capacity for exploiting parallelism. Beyond that limit, parallelism is at best worthless, and usually comes with overhead (e.g. excessive synchronisation).
- ▶ *A straightforward implementation of this idea only works if the operator is commutative, and also allows fusing a map producer!*

Applying Brent's Lemma

The idea: instead of letting the thread count depend on the input size, always launch the same number of threads, and have each thread perform an efficient sequential summation of a *chunk* of the input.

- ▶ GPUs have a maximum (hardware/problem-dependent) capacity for exploiting parallelism. Beyond that limit, parallelism is at best worthless, and usually comes with overhead (e.g. excessive synchronisation).
- ▶ *A straightforward implementation of this idea only works if the operator is commutative, and also allows fusing a map producer!*

Implementation	$n = 1000$	$n = 1000000$
Tree reduction	$77\mu s$	$363\mu s$
Block reduction	$17\mu s$	$179\mu s$
Chunked reduction	$70\mu s$	$103\mu s$

Using Atomics

The idea: GPUs have special hardware support for performing certain memory updates atomically. In CUDA, this is exposed through *atomic operations*.

```
int atomicAdd( volatile __global int *p  
               , int val)
```

- ▶ Concise parallel reduction: each thread reads an element and uses `atomicAdd()` to update the same location in memory.
- ▶ **Why is this slow for large inputs?**

Using Atomics

The idea: GPUs have special hardware support for performing certain memory updates atomically. In CUDA, this is exposed through *atomic operations*.

```
int atomicAdd( volatile __global int *p  
               , int val)
```

- ▶ Concise parallel reduction: each thread reads an element and uses `atomicAdd()` to update the same location in memory.
- ▶ **Why is this slow for large inputs?**

Implementation	$n = 1000$	$n = 1000000$
Tree reduction	$77\mu s$	$363\mu s$
Block reduction	$17\mu s$	$179\mu s$
Chunked reduction	$70\mu s$	$103\mu s$
Atomics	$8\mu s$	$1278\mu s$

Coalesced Access to Global Memory

- ▶ On NVIDIA GPUs, the hardware threads are split into WARPS, where a WARP consists of 32 threads.
- ▶ The threads in a WARP execute in lockstep—in SIMD fashion—meaning at each given cycle they all execute the same instruction.
- ▶ Coalesced access to global memory is obtained when the threads in a WARP access in their common load/store instruction consecutive words in global memory.
- ▶ A coalesced load/store instruction is serviced by one (possibly two) memory transactions, hence fast.
- ▶ Uncoalesced access may requires as many as 32 memory transactions, which are sequentially issued by the memory controller, hence may generate significant slowdown.
- ▶ Multi-threading somehow alleviates the “uncoalesced” overhead but only to some extent.

For Non-Commutative Operators (MSSP)

We denote the CUDA block size with B .


- ▶ We can also fuse the mapped function in the reduction itself.
- ▶ Have each thread reducing sequentially some CHUNK **consecutive** elements (statically known):
 - ▶ then the B per-thread results are reduced cooperatively by the threads in a block.
 - ▶ sequential (virtualization) loop on top so that the number of blocks is kept to under 1024, i.e., two-stage reduce.
- ▶ However, if done naively, this will lead to **uncoalesced** access. For **coalesced**, use shared memory as a staging buffer:
 - ▶ the threads in a block read consecutive elements from global memory and write them in shared memory;
 - ▶ then each thread processes sequentially its CHUNK consecutive elements from shared memory, which is not subject to the “coalescing” overhead.

Course Contents

Implementation of Reduce

Implementation of Scan

Parallel Exclusive Scan with Associative Operator \oplus




Two Steps:

- ▶ **Up-Sweep:** similar with reduction
- ▶ Root is replaced with neutral element.
- ▶ **Down-Sweep:**
 - ▶ the left child sends its value to parent and updates its value to that of parent.
 - ▶ the right-child value is given by \oplus applied to the left-child value and the (old) value of parent.
 - ▶ note that the right child is in fact the parent, i.e., in-place algorithm.

Scan's Work and Depth: $D(n) = \Theta(\lg n)$, $W(n) = \Theta(n)$

Warp-Level Inclusive Scan for GPUs



Input: array A of $n=2^k$ elements
of type α


$\oplus : (\alpha, \alpha) \rightarrow \alpha$ associative

Output: $B = [a_1, a_1 \oplus a_2, \dots, \oplus_{j=0}^{n-1} a_j]$

1. forall $i = 0 : n-1$ do
2. $B[i] \leftarrow A[i]$
3. endfor
4. for $d = 0$ to $k-1$ do
5. $h = 2^d$
6. forall $i = h$ to $n-1$ do
7. $B[i] \leftarrow B[i-h] \oplus B[i]$
8. endfor
9. endfor

Offers better performance because it operates in one sweep rather than two!

Warp-Level Inclusive Scan Implementation



Input: array A of $n=2^k$ elements
of type α

$\oplus : (\alpha, \alpha) \rightarrow \alpha$ associative

Output: $B = [a_1, a_1 \oplus a_2, \dots, \oplus_{j=0}^{n-1} a_j]$

1. forall $i = 0 : n-1$ do

2. $B[i] \leftarrow A[i]$

3. endfor

4. for $d = 0$ to $k-1$ do

5. $h = 2^d$

6. forall $i = h$ to $n-1$ do

7. $B[i] \leftarrow B[i-h] \oplus B[i]$

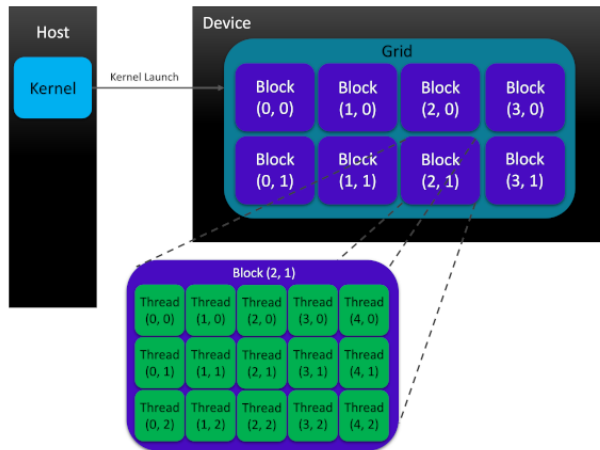
8. endfor

9. endfor

- ▶ Open file "pbbKernels.cu.h", and implement function named "scanIncWarp" (follow the instructions)
- ▶ Your $n = \text{WARP}$ and $k = \lg \text{WARP}$; Ignore the init loop;
- ▶ Unroll the for d loop (#pragma unroll);
- ▶ loop forall $i = h$ to $n-1$ is implicit (parallel threads),
- ▶ it should be replaced by a condition if $(i \geq h) \{ \dots \}$,
- ▶ except that i in condition if $(i \geq h)$ **is not the thread id**.
- ▶ Remember, you want to scan each warp, independently!

Cuda: Grid-Block Structure of Threads

Credit: pictures taken from http://education.molssi.org/gpu_programming_beginner/03-cuda-




Blocks and Grids have at most three dimensions—denoted x, y, z , with x innermost and z outermost. Their sizes are specified at kernel launch. Inside the **kernel** you may use:

- ▶ `blockDim.x`: block size in dim x
- ▶ `blockIdx.x`: current block index (in x)
- ▶ `threadIdx.x`: local index of the current thread inside its block (in dim x)
- ▶ `gridDim.x` number of blocks on dim x
- ▶ Ditto for dimensions y and z .

The global thread index in dim $q \in \{x, y, z\}$:

$\text{threadIdx.q} + \text{blockIdx.q} \cdot \text{blockDim.q}$



CUDA Scan Implementation

BLACKBOARD!

- ▶ Generic CPU skeleton in `hostSkel.cu.h`.
 1. reduce each block and publish the per-block results in a buffer `buf` (number of blocks ≤ 1024)
 2. scan in place the buffer `buf` using one CUDA block.
 3. now you can implement the core scan kernel:
 - 3.1 each thread reads and scans its `CHUNK` of consecutive elements held in register memory; (The elements are copied in a coalesced way by using shared memory as a staging buffer.)
 - 3.2 the last element (reduction) of each thread's chunk (held in registers) is published in shared memory.
 - 3.3 then the threads in the block cooperatively scan the per thread reductions.
 - 3.4 then each thread `tid` updates the elements of its chunk array with the element at position `tid-1` from step 3.3, and with the previous-block element of `buf` from step 2.
 - 3.5 the `CHUNK` per-thread results are copied from register to shared to global memory (to ensure coalesced writes)
- ▶ If N is the length of the input, this requires $2N$ coalesced reads and N coalesced writes from/to global memory.